

Real-Time OS Frameworks

Raj Rajkumar

Lecture #7

Outline of Lecture

- Approaches limiting Real-time and Non-real-time Task Interactions
 - Compliant Kernel Approach
 - Dual/Thin Kernel Approach
- Approaches that integrate Real-time and Non-real-time tasks
 - Core Kernel Approach
 - Resource Kernel Approach

Approaches to Real-Time Linux

- Approaches limiting Real-time and Non-real-time Task Interactions
 - Compliant Kernel Approach
 - LynxOS/Blue Cat Linux
 - Dual Kernel Approach
 - RTLinux/RTAI
- Approaches that integrate Real-time and Non-real-time tasks
 - Core Kernel Approach
 - Monta Vista Linux, TimeSys Linux
 - Resource Kernel Approach
 - Linux/RK

Linux Internals: Scheduling

- Schedulable Entities
 - Processes
 - Real-Time Class: `SCHED_FIFO` or `SCHED_RR`
 - Time-Sharing Class: `SCHED_OTHER`
 - Real-time processes have
 - Application-defined priority
 - Higher priority than time-sharing processes
- Non-Schedulable Entities
 - Interrupt Handlers
 - Have priorities, and can be nested
 - “Bottom Halves” & Task Queues
 - Run on schedule, ret from system call, ret from interrupt

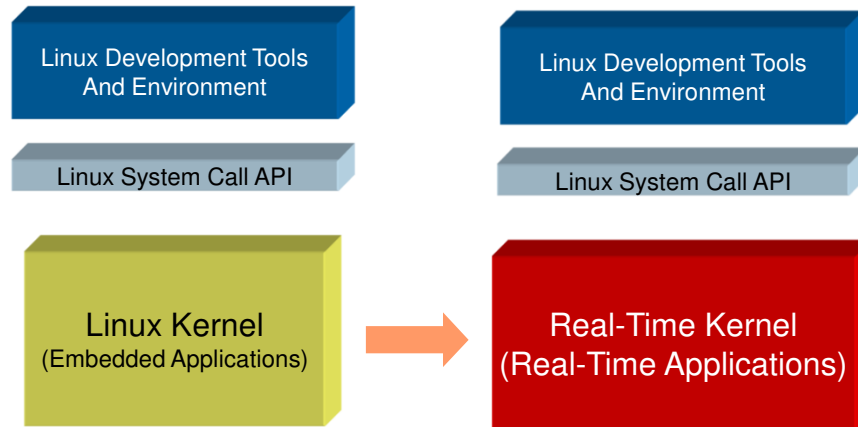
Linux and Real-Time: Traditional Problems

- **Timer Granularity**
 - Many real-time tasks are driven by timer interrupts
 - In Standard Linux, the timer was set to expire at 10 ms intervals
 - Beginning to change with usage of high-resolution timer and timestamp counters
- **Scheduler Predictability**
 - The Linux scheduler used to keep tasks in an unsorted list
 - Requires a scan of all tasks to make a scheduling decision
 - Scales poorly as number of tasks increases, and is especially poor for real-time performance
- **Various subsystems NOT designed for real-time use**
 - Network protocol stack
 - Filesystem
 - Windows manager

Approaches to Real-Time Linux

- Compliant Kernel Approach
- Dual Kernel Approach
- Core Kernel Approach
- Resource Kernel Approach

Compliant Kernel Approach



Compliant Kernel Approach

- Basic Claim
 - Linux is defined by its API and not by its internal implementation
 - The real-time kernel is a non-Linux kernel
- Implications
 - No benefits from the Linux kernel
 - Not possible to benefit from the Linux kernel evolution
 - Not possible to use Linux hardware support
 - Not (always) possible to use Linux device drivers

Compliance

- 100% Linux API
 - Support all of Linux kernel API
- Implications
 - + Any Linux application can run on real-time kernel
 - Development can be done on a Linux host, with a rich set of host tools for development
 - + All Linux libraries are trivially available to run on a real-time kernel
 - Third-party software
 - Achieving 100% Linux API is non-trivial
 - Consider the amount of effort put into Linux kernel development

Approaches to Real-Time Linux

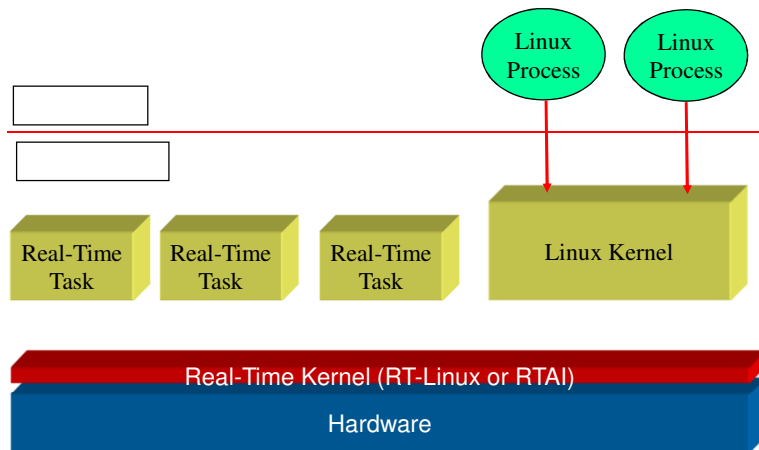
➤ Compliant Kernel Approach

➤ Dual Kernel Approach

➤ Core Kernel Approach

➤ Resource Kernel Approach

The Thin Kernel Approach



- Real-time tasks do NOT use the Linux API or Linux facilities
- Failure in any real-time task crashes the entire system

Approaches to Real-Time Linux

- Compliant Kernel Approach
- Dual Kernel Approach
- Core Kernel Approach
- Resource Kernel Approach

Core Kernel Approach

- **Basic Ideas**
 - Make the kernel more suitable for real-time
 - Ensure that the impact of changes is localized so that
 - Kernel upgrades can be easily incorporated
 - Kernel reliability and scalability is not compromised
- **Mechanisms**
 - Static Configuration
 - Can be configured at compile time
 - Dynamic Configuration
 - Using loadable kernel modules

Core Kernel Approach

- Allows the use of most, if not all, existing Linux primitives, applications, and tools.
 - Need to avoid primitives that can take extended time in the kernel
- Allows the use of most existing device drivers written to support Linux.
 - Need to avoid poorly written drivers that unfairly hog system resources
- **Robustness and Reliability**
 - Core kernel modifications can affect robustness, but source is available and extensive testing can be done.

Approaches to Real-Time Linux

- Compliant Kernel Approach
- Dual Kernel Approach
- Core Kernel Approach
- Resource Kernel Approach

Resource Kernel

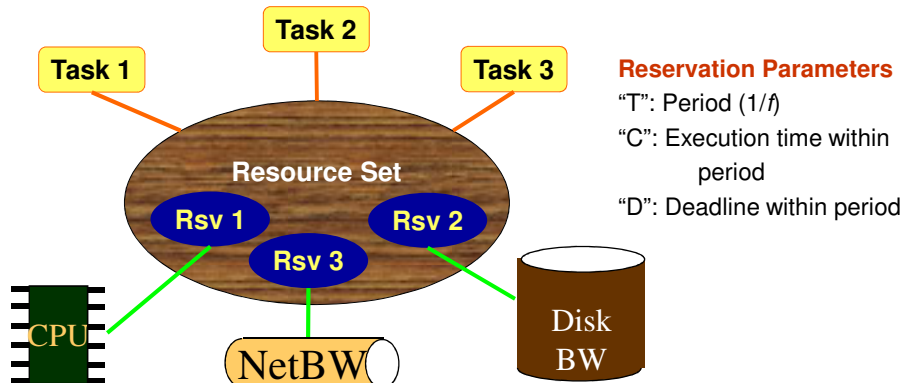
- A Kernel that provides to Applications Timely, Guaranteed, and Enforced access to System Resources
- Allows Applications to specify only their Resource Demands, leaving the Kernel to satisfy those Demands using hidden management schemes

Protection in Resource Kernels

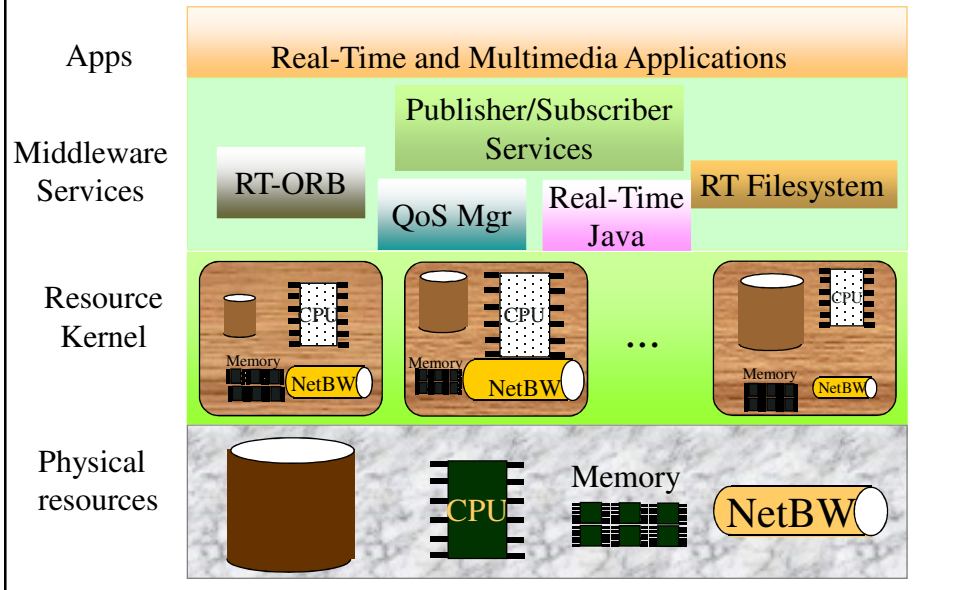
- Each application (or a group of collaborating applications) operates in a “virtual machine”:
 - a machine which consists of a well-defined and guaranteed portion of system resources
 - CPU capacity, disk bandwidth, network bandwidth, and memory resource
- Multiple virtual machines can run simultaneously on the same physical machine
 - guarantees available to each resource set is valid despite the presence of other (potentially mis-behaving) applications using other resource sets

Resource Kernel


- A Kernel that provides to applications **Timely, Guaranteed, and Enforced access to System Resources**
- Allows Applications to specify only their Resource Demands
 - leaving the Kernel to satisfy those Demands using hidden management schemes



“Resource Kernel” Architecture



Linux Resource Kernel Architecture



Reserves and Resource Sets


- **Reserve**
 - A Share of a Single Resource
 - Temporal Reserves
 - Parameters declare Portion and Timeframe of Resource Usage
 - E.g., CPU time, link bandwidth, disk bandwidth
 - Spatial Reserves
 - Amount of space
 - E.g., memory pages, network buffers
- **Resource Set**
 - A set of resource reserves
 - Zero, one or more processes can be bound to a resource set.

Linux/RK Abstractions

Linux/RK supports several abstractions and primitives for real-time scheduling of processes with real-time and QoS requirements:


- **Resource reservations with latency guarantees**
 - CPU cycles
 - Network bandwidth
 - Disk bandwidth
- Support for **periodic tasks**.
- Support for 256 real-time fixed-priority levels.
- **High-resolution timers and clocks**.
- **Bounding of priority inversion** during synchronization operations.
- Wiring down of memory pages.

Reservation Types




Hard Reservations

(guarantees with **No** extras even if resource is idle)



Firm Reservations

(guarantees with Extras only if no non-real-time)



Soft Reservations

(guarantees with extras)

Carnegie Mellon

18-648: Embedded Real-Time Systems

Electrical & Computer
ENGINEERING

Performance Overhead of Hierarchical Reservations

- **Linear with the height of the hierarchy**
 - reservation replenishment and enforcement
- **Constant**
 - admission control (only *local* schedulability analysis)
 - scheduling (internal priority mapping and disabling/re-enabling of process eligibility to be scheduled).
- **Constraints**
 - reservation period must be greater than twice the parent's reservation period
- **Hidden overhead**
 - Higher degree of interrupts, because of replenishment, enforcement timers going off more frequently.



Carnegie Mellon

18-648: Embedded Real-Time Systems

Electrical & Computer
ENGINEERING

Degrees of Temporal Isolation

- Different degrees of temporal isolation in the presence of resource-sharing
 - **Strict Isolation:** the timing behavior of an application is not affected by the timing misbehavior of *any* other application
 - RK applications in the absence of logical sharing of resources
 - **Non-Strict Isolation:** traditional priority-driven systems
 - **Weak Isolation:** timing behavior is not affected by the timing misbehavior of applications with which no logical resources are shared.



Resource-Sharing Protocols in RK

- **Analogues to Priority Inheritance and Priority Ceiling Protocols** in Resource kernels
- Temporal isolation can only be weak
 - under logical resource-sharing using mutexes and client-server architectures
 - timeout and restart schemes may need to be applied.

Priority Ceiling Protocol Equivalents

- **Single-Reserve PCP**

- Assign one reservation to the logical resource execution
- Very pessimistic allocation is required to maintain PCP semantics

- **Multi-Reserve PCP**

- Has the same schedulability analysis as traditional PCP
- Requires special support in RK
- Can be applied to client-server models
 - Pass client's reserve to server along with request - charges

Priority Inheritance Protocol Equivalents

- **Priority Inheritance (PI)**

- Server runs at the priority of the highest priority client waiting for server.


- **Priority Inheritance with Priority Inversion Enforcement (PIPIE)**

- Enforce the duration of priority inversion encountered by any task
 - Can never exceed the amount specified at admission control
 - Need to track multiple tasks' priority inversions simultaneously


- **Reserve Inheritance (RI)**

- Server usage is charged to client's reservation + inherit the highest priority of any client waiting for server

- **PIPIE + RI**



Network Processing Architecture



Problems

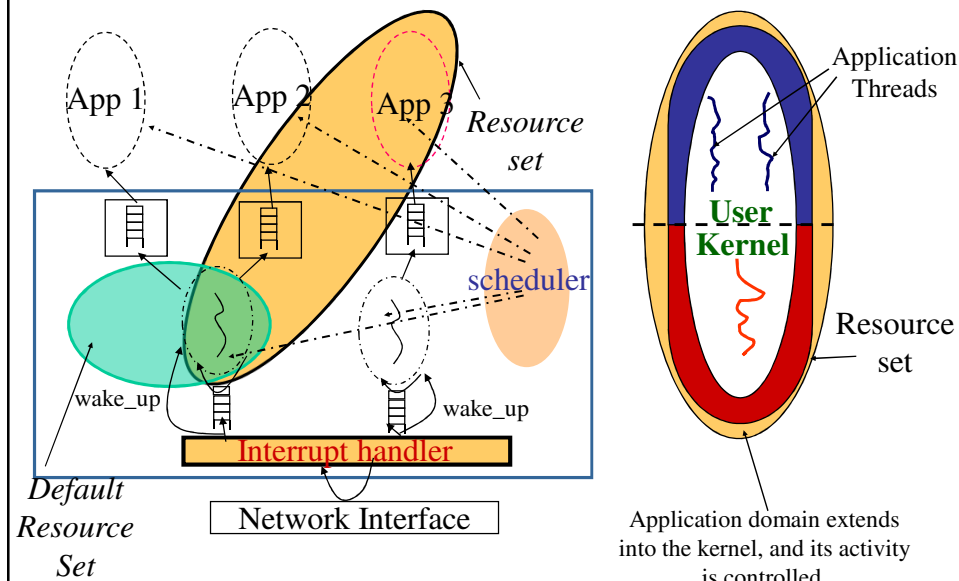
- Non-preemptive System calls
 - Priority Inversion
- “Eager” Receiver Processing
 - Interrupt-driven
 - Priority(capturing packet) > Priority(protocol processing) > Priority(application processing)
- Lack of effective load shedding
- Lack of traffic separation
- Inappropriate resource accounting

Carnegie Mellon

18-648: Embedded Real-Time Systems

Electrical & Computer ENGINEERING

Threaded Network Processing



Carnegie Mellon

18-648: Embedded Real-Time Systems

Electrical & Computer ENGINEERING

Network and CPU Service Guarantees

- Reduction in non-preemptibility
- Control of receiver overload (*receive-livelock*)
- Prevention of scheduling disruption
- Separation of individual flows and proper resource accounting
- Packet scheduling for QoS (Quality of Service) guarantees

Summary

- OS Approaches limiting Real-time and Non-real-time Task Interactions
 - Compliant Kernel Approach
 - Thin Kernel Approach
- OS Approaches that integrate Real-time and Non-real-time tasks
 - Core Kernel Approach
 - Resource Kernel Approach