

# Course Summary

Raj Rajkumar

Lecture #24

## Final Exam

- Final Exam on Sunday, December 17, 2017
  - 8:30 - 11:30 AM
  - Baker Hall A51

## COURSE SUMMARY

## Goals of the Course

- Understand the scientific principles and concepts behind embedded real-time systems, and
- Obtain hands-on experience in programming embedded real-time systems (in the form of smartphones)
  - Understand the “big ideas” in embedded real-time systems.
  - Obtain direct hands-on experience.
  - Understand basic real-time resource management theory.
  - Understand the basics of embedded real-time system application concepts such as signal processing and feedback control.

## Introduction to Embedded Systems

- What is an embedded system?
  - It is a device that needs to accomplish a particular set of tasks. Over time though, the distinction between an embedded system and a computer is becoming fuzzier but is still distinct in many cases.
- Attributes:
  - Reactive
  - Real-Time
- Typical Constraints:
  - Small Size, Low Weight
  - Low Battery, Harsh Environment
  - Safety-critical operation
  - Extreme Cost Sensitivity
  - Timing Constraints: Hard deadlines, Soft deadlines.

## OS Basics

- Process Management:
  - Process States
  - Process Control Block
  - Threads vs Processes
  - Process creation: fork, exec
- RTOS: includes both logical and temporal correctness
- Resource Availability: Abundant, Insufficient, Sufficient but Scarce
- Tasks: Periodic, Sporadic and Aperiodic
- Scheduling: Static and Dynamic
- Rate-Monotonic Analysis

## Review Question

How many times will the following program print “Hello World!”?

```
#include <sys/types.h>
#include <unistd.h>
#include <stdio.h>
#include <stdlib.h>

int n = <some positive integer>;

int main() {
    int i;

    for (i = 0; i < n; i++) {
        fork();
        fork();
        printf("Hello World!\n");
    }
    exit(0);
}
```

## Android

- From Google
- Uses Java
- Linux Kernel 2.6
- Dalvik Virtual Machine
- Provides Application Framework
- Application-level Power Management done through “wakelocks”
- Provides debugging interface through ADB



## Uniprocessor Scheduling

- **Earliest Deadline First:**
  - Dynamic priority, Higher Priority to task with earlier deadline
  - Task-set is schedulable if  $U \leq 1$
  - If deadlines are shorter than periods, necessary condition for schedulability under EDF is an open problem.
  - If  $U > 1$ , which task will miss its deadline is unpredictable.
- **Rate-Monotonic Scheduling:**
  - Higher priority to task with shorter period
  - RMS leads to a feasible schedule if  $U \leq n(2^{1/n} - 1)$
  - The above bound is sufficient but not necessary
- **Critical Instant:** The instant or relative phasing at which a task arrival encounters its worst-case response time.
- **Critical Zone:** The duration between the critical instant of a task instance and its completion.

## Review Question

Consider the following task-set with the following  $\tau = (C, T=D)$  parameters:  $\{(5, 10), (5, 20), (5, 40), (5, 80)\}$ . Draw the time-lines for this task-set at least until  $t=120$ . Assume that all tasks are released at  $t=0$ . When deadlines or priorities are equal, break ties on a FCFS (first-come-first-served) basis.

EDF



## Response Time Test

Consider the task set  $\{C, T=D\} = \{(20, 100), (90, 150), (60, 300)\}$ .

$$a_{k+1}^i = C_i + \sum_{j=1}^{i-1} \left\lceil \frac{a_k^i}{T_j} \right\rceil C_j \quad \text{where } a_0^i = \sum_{j=1}^i C_j$$

Test terminates when  $a_{k+1}^i = a_k^i$

For task  $\tau_1$

$$a_0^1 = 20$$

$$a_0^2 = C_1 + C_2 = 20 + 90 = 110$$

For task  $\tau_2$

$$a_1^2 = C_2 + \left\lceil \frac{110}{100} \right\rceil * 20 = 130$$

$$a_2^2 = 90 + \left\lceil \frac{130}{100} \right\rceil * 20 = 130$$

## Response Time test

For task  $\tau_3$

$$a_0^3 = C_1 + C_2 + C_3 = 90 + 60 + 20 = 170$$

$$a_1^3 = C_3 + \left\lceil \frac{170}{100} \right\rceil * 20 + \left\lceil \frac{170}{150} \right\rceil * 90 = 280$$

$$a_2^3 = 60 + \left\lceil \frac{280}{100} \right\rceil * 20 + \left\lceil \frac{280}{150} \right\rceil * 90 = 300$$

$$a_3^3 = 60 + \left\lceil \frac{300}{100} \right\rceil * 20 + \left\lceil \frac{300}{150} \right\rceil * 90 = 300$$

## Real Time Synchronization

- **Priority Inversion:** When lower-priority tasks cause higher-priority tasks to wait.
- **Critical Section:** The duration of a task using a shared resource.
  - Can potentially lead to unbounded delay of a task execution.
- Sources: Synchronization and mutual exclusion, Non-preemptible regions of code.
- **Basic Priority Inheritance Protocol**
  - Let the lower priority task use the highest priority of the higher priority task that it blocks.
  - There will be no deadlocks IF there are no nested locks, or application level deadlock avoidance scheme.
- **Priority Ceiling Protocol**
  - A priority ceiling is assigned to each mutex, which is equal to the highest priority task that may use this mutex.
  - A task can lock its mutex if and only if its priority is higher than the priority ceilings of all mutexes locked by other tasks.
  - If a task is blocked by a lower priority task, the lower priority task inherits its priority.
- **Highest Locker's Priority Protocol**
  - Execute critical section immediately at priority = priority ceiling of critical section

## Review Question

- Consider the task set  $\tau_1 = (40, 100, 100)$ ,  $\tau_2 = (40, 200, 200)$ ,  $\tau_3 = (150, 600, 600)$ , which specifies the (C,T,D) triplet for each task  $\tau_i$ . These tasks access one or more critical sections guarded by 3 mutexes  $M_1$ ,  $M_2$  or  $M_3$ .
  - Within its execution time,  $\tau_1$  accesses a critical section guarded by  $M_1$  for 15 ms; and after normal execution, it enters a critical section guarded by  $M_2$  for 18 ms.
  - During its execution,  $\tau_2$  accesses a critical section guarded by  $M_2$  for 13 ms, and after normal execution, it enters a critical section guarded by  $M_3$  for 10 ms.
  - During its execution,  $\tau_3$  accesses a critical section guarded by  $M_1$  for 20 ms and after normal execution, it enters a critical section guarded by  $M_3$  for 14 ms.
- Blocking time when using:
  - PIP? PCP? Highest locker? Non-preemption protocol?

## POSIX pthreads

- POSIX is a standard set of interfaces for operating systems, standardized by the IEEE.
- Pthreads is the set of POSIX interfaces for threads
- Support interfaces for
  - Threads management, thread attributes
  - Mutex management, mutex attributes
  - Condition variable management, condition variable attributes
- Real-time scheduling policies to support basic priority inheritance and priority ceiling protocols are included.

## RTOS Options

- **OS Approaches limiting Real-time and Non-Real-time Task interactions**
  - **Compliant Kernel Approach:** 100% Linux API, Any application can run on real-time kernel
  - **Thin Kernel Approach:** RT-Linux or RTAI
- **OS Approached that integrate Real-time and Non Real-time tasks**
  - **Core Kernel Approach:** Make the kernel suitable for real-time while changes to the kernel are localized. Allows the use of most if not all existing Linux primitives, applications and tools.
  - **Resource Kernel Approach :** Provides applications Timely, Guaranteed, and Enforces access to System Resources. Applications need to specify their resource demands only.




## Deadline-Monotonic Scheduling

- Assign fixed priority based on  $D$  (and not  $T$ )
  - Shorter the relative deadline, higher the priority.
- When  $D=T$ , RMS and DMS are one and the same.
- When  $D > T$ , neither RMS nor DMS is the optimal fixed-priority scheduler.
- DMS is optimal among fixed-priority preemptive schedulers for periodic real-time tasks when  $D \leq T$ .
- RMS is optimal among fixed-priority preemptive schedulers for periodic real-time tasks when  $D = T$ .
- The general set of principles for analyzing fixed-priority preemptive scheduling policies is called RMA (rate-monotonic analysis).

## Review Question

Task	$C_i$	$T_i$	$D_i$	Priority
$\tau_1$	1	4	4	
$\tau_2$	1	5	3	
$\tau_3$	1	6	6	
$\tau_4$	1	10	5	



## Periodic Servers

- **Polling Server:**
  - Simple
  - Commonly used
  - WCRT can be long.
- **Deferrable Server:**
  - Improves response time
  - Cannot be generalized to multiple instances at different priority levels.
- **Sporadic Server:**
  - Improves upon the deferrable server and is very generalizable
  - Higher run-time complexity

## Review Question

Draw the execution timeline for the tasks  $\tau_1$ : ( $C = 2$ ,  $D = T = 8$ ),  $\tau_2$ : (2,10) and a deferrable server task  $\tau_s$ : (2,6) till  $t=24$  where:

- All the tasks are scheduled using RMS.
- The arrival and computation times of aperiodic tasks associated with the server are:

Arrival Time	Computation Time
5	2
9	1
12	2
16	1

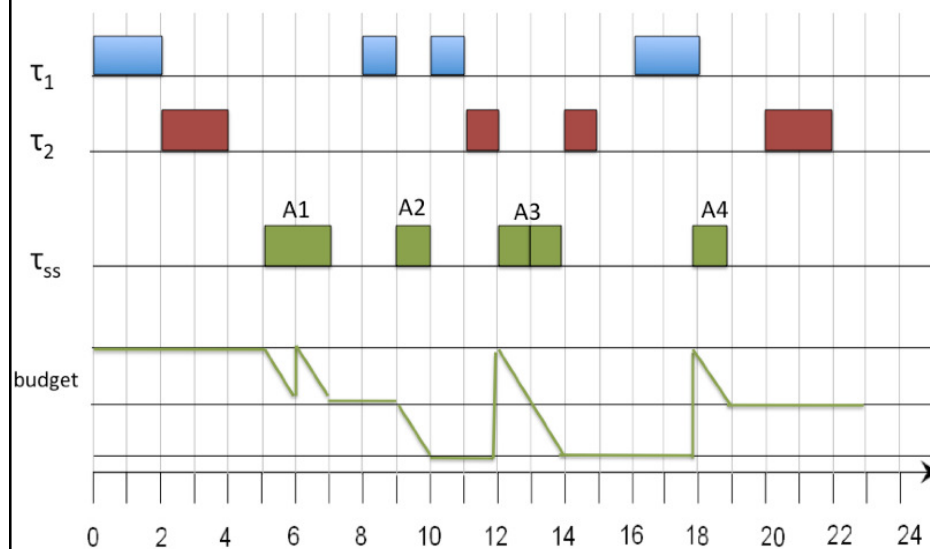
## Review Question

Draw the execution timeline for the tasks  $\tau_1$ : ( $C = 2$ ,  $D = T = 8$ ),  $\tau_2$ : ( $2$ ,  $10$ ) and a sporadic server task  $\tau_s(2,6)$  till  $t=24$  where:

- All the tasks are scheduled using RMS.
- The arrival and computation times of the aperiodic task are given below:

Arrival Time	Computation Time
5	2
9	1
12	2
16	1

## Review Question



## Power Management

- Power is proportional to  $V^2f$
- Scaling Techniques:
  - DFS (dynamic frequency scaling)
  - DVS (dynamic voltage scaling)
  - DVFS (dynamic voltage and frequency scaling)
- Static Voltage Scaling Algorithm
  - **Sys-Clock**: Optimal system-wide clock frequency assignment
  - **PM-Clock**: Task-specific clock frequency assignment
    - Longer the period, lower the clock frequency
- Static Voltage Scaling with EDF
- Cycle-conserving EDF
- Harmonized scheduling
  - **Rate-Harmonized Scheduling**
  - **Energy-Saving Rate-Harmonized Scheduling**

## Review Question

You are given two tasks  $\tau_1$  and  $\tau_2$  defined as  $\tau_1 = (C = 3, T = 6, D = 5)$  and  $\tau_2 = (1, 30, 30)$ . Assume that the maximum frequency you can run the processor is  $f_{max}$ .

What is the Sys-Clock frequency?

Ans:

What are the clock-frequency assignments under PM-Clock?

Ans:

0.6

Task 1: 0.6 Task 2: 0.2

## Review Question

Consider the following taskset scheduled on a micro-controller with a sleep duration  $C_{\text{sleep}}$  of 8.


$$\tau_1 = \{C_1 = 4, T_1 = 16\}$$

$$\tau_2 = \{C_2 = 2, T_2 = 27\}$$

$$\tau_3 = \{C_3 = 2, T_3 = 40\}$$

Simulate the timeline up to 40 for the given tasks under  
RHS with a harmonizing period of 8  
ES-RHS with a harmonizing period of 16

## Answer to Review Question



## Multi-processor Scheduling

- *Global scheduling* and *partitioned scheduling* are two approaches to multiprocessor scheduling.
- Timing behavior under task scheduling strategies can be brittle. Small changes can have big (unexpected) consequences.
- **Bin-Packing Heuristics** for partitioned scheduling.
  - The problem is known to be NP-complete.
  - Very efficient near-optimal heuristics exist.
  - Problems with bin-packing?
- **Task Splitting** to circumvent 50% bound
  - Highest-priority task splitting
  - Period transformation
  - Task splitting for harmonics

## Review Question

Consider the following task-set with 6 tasks (Assume  $D_i = T_i$ ). Allocate the tasks to processors when tasks cannot be split. Mission: Minimize bin count.

$$\tau_1: \{C_1 = 4, T_1 = 10\}$$

$$\tau_2: \{C_2 = 6, T_2 = 14\}$$

$$\tau_3: \{C_3 = 4, T_3 = 16\}$$


$$\tau_4: \{C_4 = 4, T_4 = 16\}$$

$$\tau_5: \{C_5 = 3, T_5 = 10\}$$

$$\tau_6: \{C_6 = 3, T_6 = 10\}$$

## Review Question

First Fit Decreasing




Carnegie Mellon

18-648: Embedded Real-Time Systems

Electrical & Computer  
ENGINEERING

## Review Question

Best-Fit Decreasing



Carnegie Mellon

18-648: Embedded Real-Time Systems

Electrical & Computer  
ENGINEERING

## Review Question


Consider the following task-set with 3 tasks (Assume  $D_i = T_i$ ).

$$\tau_1 = \{C_1 = 3, T_1 = 5\}$$

$$\tau_2 = \{C_2 = 6, T_2 = 10\}$$

$$\tau_3 = \{C_3 = 12, T_3 = 20\}$$

BFD without task-splitting:




Carnegie Mellon

18-648: Embedded Real-Time Systems

Electrical & Computer  
ENGINEERING

## Review Question

BFD with task-splitting



Carnegie Mellon

18-648: Embedded Real-Time Systems

Electrical & Computer  
ENGINEERING



## Real-Time Communications

- Timely delivery may be deemed more desirable than reliable delivery.
- There are many causes of delay:
  - Queuing delay at sender, receiver, network and the propagation delay.
- Throughput and delay depend on the capacity of the link.
- The requirements of throughput and delay are application dependent.
- The buffering depends on the jitter.
- **Controller Area Network (CAN):**
  - Connections wired together as a logical AND function.
  - Distributed Priority-based Arbitration.
- **FlexRay**
  - Delivers deterministic, fault-tolerant and high-speed bus system.

## Feedback Control

- Feedback control
  - Stability
  - Instability
  - Marginal stability
- Feedback controllers
  - Proportional control
  - Proportional + derivative control
  - Proportional + derivative + integral control

## Signal Processing

- Source of deterministic errors and random noises
- Basics of signal spectrum
  - Nyquist sampling
    - the sampling rate must be at least two times the bandwidth, the highest frequency in the signal, to avoid aliasing. This is known as the **Nyquist Rate**.
  - Fourier transform
- Basic filters
  - If there is no aliasing, then the signal can be recovered perfectly *in theory* using ideal **low-pass filters**.

## Final Exam

- An expanded, sometimes more in-depth, version of the 5 quizzes to date.
  - True/False
  - Short questions
  - Long questions with multiple sub-parts
- Bring a calculator, pencil and eraser.
- Do *not* get stuck on any question.
  - Skip to later/easier ones and return to tougher ones later
- Time limits will be strictly enforced.

**BEST WISHES**

**Carnegie Mellon**

*18-648: Embedded Real-Time Systems*

 **Electrical & Computer  
ENGINEERING**