

Name: _____ **SOLUTIONS** _____ (Bonus: 1 point)

18-648: Embedded Real-Time Systems

Quiz #1

Fall 2017

60 minutes

Instructions

1. This is a **CLOSED-BOOK/CLOSED-NOTES** quiz.
2. Show all relevant work.
3. Partial credit may be given for some questions.
4. The use of a calculator is allowed.
5. The time limit will be *strictly* enforced.
6. Watch the screen for any clarifications.

For Graders' Use Only

Bonus: _____ / 1

1. _____ / 15

2. _____ / 4

3. _____ / 19

4. _____ / 12

TOTAL. _____ / 50

Question 1. Fast Food.

(a) True/False Questions (12 points)

- i. **FALSE** A hard real-time system can only use hardware, while a soft real-time system can also use software.
- ii. **TRUE** Given a periodic taskset, if any fixed-priority preemptive scheduling policy can render that taskset schedulable, then this taskset is also schedulable using the rate-monotonic scheduling policy.
- iii. **TRUE** Under rate-monotonic scheduling, when the relative deadline is at the end of a task period, the schedulability of a task can be tested by checking its first deadline when all tasks arrive together.
- iv. **FALSE** In a real-time system, processing must be done as fast as possible.
- v. **TRUE** If every thread, that holds two mutexes $m1$ and $m2$ simultaneously, locks them in the same order, the system will never deadlock.
- vi. **TRUE** Threads of the same process share their data and code segments.
- vii. **TRUE** Using the Priority Ceiling Protocol (PCP) will bound the blocking time due to priority inversion.
- viii. **TRUE** Only a limited number of threads can be created within a process.
- ix. **FALSE** With a priority-based preemptive scheduler, a ready task scheduled to run will run until its completion, after which the highest-priority ready task on the ready queue will be run.

- x. **FALSE** Fixed-priority preemptive schedulers have a higher schedulable utilization bound than dynamic-priority preemptive schedulers.
- xi. **FALSE** Earliest deadline first (EDF) is a fixed-priority preemptive-scheduling scheme.
- xii. **TRUE** Consider a fixed-priority scheduler which assigns random (fixed) priority values to periodic real-time tasks. There will be *never* be any priority assignments on some task sets for which this scheduler would find a feasible schedule but the rate-monotonic scheduler will not.

(b) Circle one or more answers for the questions below. (3 points)

- I. User-space applications access kernel functions through:
- a. A shell.
 - ☒ b. System calls.
 - c. Shared memory.
 - d. All of the above.
 - e. None of the above.
- II. Which of the following protocols disallow a mutex M from being locked even if M is currently unlocked?
- a. Basic Priority Inheritance Protocol (PIP)
 - b. Highest Locker Priority Protocol (HLP)
 - ☒ c. Priority Ceiling Protocol (PCP)
 - d. Non-Preemption Protocol (NPP)
- III. The schedulability of a periodic taskset using rate-monotonic scheduling can potentially be determined using
- ☒ a. A utilization bound test
 - ☒ b. A harmonicity test
 - ☒ c. Completion-time test

Question 2. At a fork? Take it. (4 points)

Given the following code, what is/are the possible output(s)? Assume that fork() succeeds.

```
count = 2;
if (pid = fork()) {
    count += 6;
    printf("%d\n", count);
    count += 1;
    return 0;
}
```

```
count += 2;
printf("%d\n", count);
return 0;
```

Possible Output #1:

8
4

Possible Output #2:

4
8

Note: In reality, printf() itself is not atomic and one could potentially also have outputs like "84\n\n" and "48\n\n".

Question 3. Schedules, schedules (19 points)

Consider a periodic taskset with 3 tasks. Each task τ_i is characterized by $\{C_i, T_i\}$, its worst-case execution time and period parameters respectively. The relative deadline of each task is the same as its period.

$$\tau_1 = \{7, 10\}$$

$$\tau_2 = \{3, 15\}$$

$$\tau_3 = \{3, 30\}$$

- (a) Determine whether the above taskset is schedulable using Earliest Deadline First scheduling algorithm. (3 points)

$$U_{\text{total}} = (7/10) + (3/15) + (3/30) = 0.7 + 0.2 + 0.1 = 1.0$$

Since $U_{\text{total}} \leq 1.0$, the taskset is schedulable under the EDF policy.

- (b) Determine whether the above taskset is schedulable using the Rate-Monotonic Scheduling policy. The RMS bound is $n(2^{1/n} - 1)$. For $n = 2$, $U_b = 0.828$. For $n = 3$, $U_b = 0.78$. Use the response-time test if appropriate. The response-time test for task i is given by:

$$a_{k+1}^i = C_i + \sum_{j=1}^{i-1} \left\lceil \frac{a_k^i}{T_j} \right\rceil C_j \quad \text{where } a_0^i = \sum_{j=1}^i C_j$$

Test terminates when $a_{k+1}^i = a_k^i$

(9 points)

Applying the **utilization test**,

For τ_1 alone, $U_1 = 0.7 < U_b(1) = 1.0$ so τ_1 is schedulable.

For τ_2 , $U_1 + U_2 = 0.7 + 0.2 = 0.9 > U_b(2) = 0.828$ – so test fails. Perform the Response Time (RT) test to check if τ_2 is schedulable.

RT test for τ_2 ($i=2$)

$$A_0^2 = C_1 + C_2 = 7 + 3 = 10$$

$$A_1^2 = C_2 + 7 \cdot \text{ceiling}(10/10) = 3 + 7 \cdot 1 = 10$$

Since $A_0^2 = A_1^2 = 10$, the RT test converges at time $t = 10 < T_2 (=15)$, so τ_2 is schedulable.

For τ_3 , $U_1 + U_2 + U_3 = 1.0 > U_b = 0.78$ fails. Perform RT test to check if schedulable.

RT test for τ_3 ($i=3$)

$$A_0^3 = C_1 + C_2 + C_3 = 7 + 3 + 3 = 13$$

$$A_1^3 = C_3 + 7 \cdot \text{ceiling}(13/10) + 3 \cdot \text{ceiling}(13/15) = 3 + 7 \cdot 2 + 3 \cdot 1 = 20$$

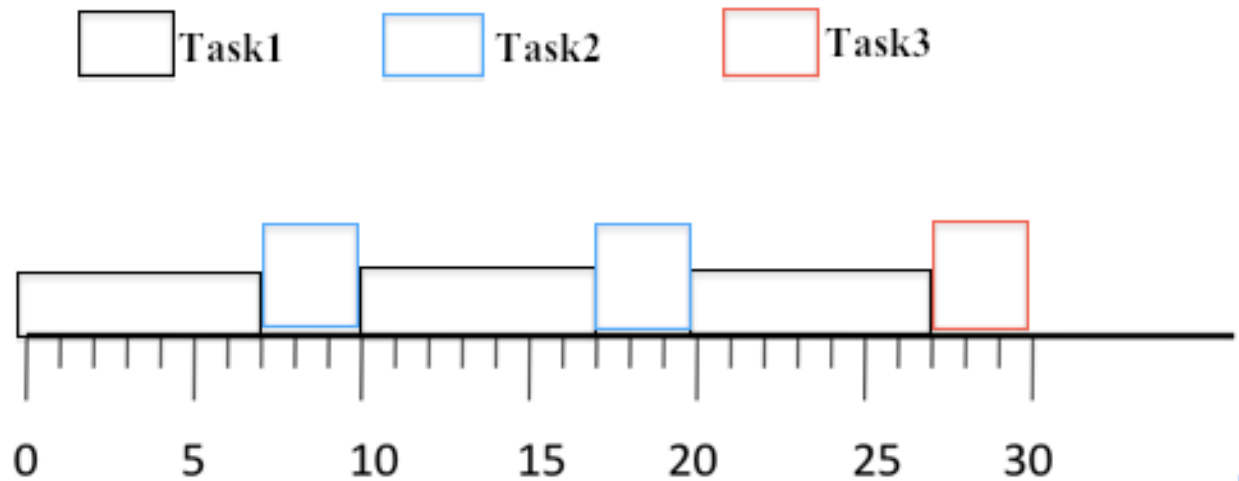
$$A_2^3 = C_3 + 7 \cdot \text{ceiling}(20/10) + 3 \cdot \text{ceiling}(20/15) = 3 + 7 \cdot 2 + 3 \cdot 2 = 23$$

$$A_3^3 = C_3 + 7 \cdot \text{ceiling}(23/10) + 3 \cdot \text{ceiling}(23/15) = 3 + 7 \cdot 3 + 3 \cdot 2 = 30$$

$$A_4^3 = C_3 + 7 \cdot \text{ceiling}(30/10) + 3 \cdot \text{ceiling}(30/15) = 3 + 7 \cdot 3 + 3 \cdot 2 = 30$$

Since $A^3_3 = A^3_4 = 30$, test converges at time $t=30 \leq T_3 (=30)$, and hence τ_3 is schedulable.

(c) Validate your result of (b) by drawing the critical zone. **(5 points)**



- The longest response time for task τ_1 is 7 ($<$ deadline/period of 10).
- The longest response time for task τ_2 is 10 ($<$ deadline/period of 15)
- The longest response time for task τ_3 is 30 (\leq deadline/period of 30).

These are consistent with our response time tests as well.

Note: The RT test mathematically computes the above manual sequence of execution to determine the longest response time for a task.

(d) Is the taskset harmonic? **(2 points)**

No. 10 and 30 are harmonic, as are 15 and 30. But 10 and 15 are not harmonic with respect to each other.

Note: The taskset has 100% utilization but is not harmonic. However, RMS is still able to schedule this taskset at 100% utilization. All harmonic tasksets will yield 100% schedulable utilization under RMS, but not all non-harmonic tasksets can do the same. The least upper utilization bound of $\ln 2$ is an example of the latter. However, the fact that some non-harmonic tasksets can reach 100% schedulable utilization for the simple RMS policy testifies to its practical effectiveness.

Question 4. Time to Block. (12 points)

A real-time taskset has 4 tasks τ_1 , τ_2 , τ_3 and τ_4 in **decreasing** order of (fixed) priority. They share some logical resources like critical sections and global variables that are protected using mutexes.

- Task τ_1 accesses a mutex M_1 for 5 *ms*.
- Task τ_2 accesses two mutexes M_2 and M_3 in non-nested fashion for 8 *ms* and 6 *ms* respectively.
- Task τ_3 accesses three mutexes M_1 , M_2 and M_4 in non-nested fashion for 10 *ms*, 12 *ms* and 14 *ms* respectively.
- Task τ_4 accesses mutexes M_3 and M_4 in non-nested fashion for 13 *ms* and 15 *ms* respectively.

Note: There are no nested mutex accesses of any kind: one mutex is locked and then released, *before* another mutex is locked.

Please fill in the table below.

Blocking Term	Under BIP	Under PCP	Under HLP	Under NPP
B_1	10	10	10	$\max(8,6,10,12,14,13,15) = 15$
B_2	$\max(10,12) + 13 = 25$	$\max(10,12,13) = 13$	$\max(10,12,13) = 13$	$\max(10,12,14,13,15) = 15$
B_3	$\max(13,15)=15$	$\max(13,15)=15$	$\max(13,15)=15$	$\max(13,15)=15$
B_4	0	0	0	0

Show work below.

Steps:

1. Determine the priority ceiling for each mutex.
2. Under each protocol, for each task, determine the set of lower-priority critical sections that can cause priority inversion to the task: this set consists of the critical sections guarded by mutexes with a priority ceiling greater than or equal to the task priority.
3. Then, apply the properties of the protocol to determine the worst-case blocking time of a task.

Priority Ceilings of Mutexes:

Priority ceiling of M_1 = priority of τ_1

Priority ceiling of M_2 = priority of τ_2

Priority ceiling of M_3 = priority of τ_2

Priority ceiling of M_4 = priority of τ_3

Under PIP:

For τ_1 , the set of lower-priority critical sections to consider: the M_1 critical section of task τ_3

For τ_2 , the set of lower-priority critical sections to consider: the critical sections of tasks τ_3 and τ_4 that access M_1 , M_2 and M_3 . Each lower-priority task can contribute only one (outermost) critical section and each mutex can contribute only one critical section. Hence, we have $B_2 = \max(10, 12)$ from task 2 + 13 from task 3.

For τ_3 , the set of lower-priority critical sections to consider: both the critical sections of task τ_4 . But a lower-priority task can contribute only one (outermost) critical section. Hence, we have $B_3 = \max(13, 15)$.

Under PCP and HLP: (these have the same blocking times for tasks when tasks do not suspend within a critical section):

For τ_1 , the set of lower-priority critical sections to consider: the M_1 critical section of task τ_3 .

For τ_2 , the set of lower-priority critical sections to consider: the critical sections of tasks τ_3 and τ_4 that access M_1 , M_2 and M_3 . Only one of these critical sections can cause priority inversion. Hence, we have $B_2 = \max(10, 12, 13)$.

For τ_3 , the set of lower-priority critical sections to consider: both the critical sections of task τ_4 . Only one of these critical sections can cause priority inversion. Hence, we have $B_3 = \max(13, 15)$.

For NPP:

Under NPP, any lower-priority critical section can cause priority inversion but at most one can. Hence, B_i is given by the longest critical section of all its lower-priority tasks.

Another equivalent alternative is to consider the priority ceilings of all mutexes to be the highest priority in the system. Hence, all lower-priority critical sections are eligible to cause priority inversion to a task but only one of them can. For the worst case, we pick the longest.

Question 4. Time to Block. (12 points) - original question on paper without fix during exam

A real-time taskset has 4 tasks τ_1 , τ_2 , τ_3 and τ_4 in **increasing** order of (fixed) priority. They share some logical resources like critical sections and global variables that are protected using mutexes.

- Task τ_1 accesses a mutex M_1 for 5 *ms*.
- Task τ_2 accesses two mutexes M_2 and M_3 in non-nested fashion for 8 *ms* and 6 *ms* respectively.
- Task τ_3 accesses three mutexes M_1 , M_2 and M_4 in non-nested fashion for 10 *ms*, 12 *ms* and 14 *ms* respectively.
- Task τ_4 accesses mutexes M_3 and M_4 in non-nested fashion for 13 *ms* and 15 *ms* respectively.

Note: There are no nested mutex accesses of any kind: one mutex is locked and then released, *before* another mutex is locked.

Please fill in the table below.

Blocking Term	Under BIP	Under PCP	Under HLP	Under NPP
B_1	0	0	0	0
B_2	5	5	5	5
B_3	$\max(8,6) + 5 = 13$	$\max(5,8,6)=8$	$\max(5,8,6)=8$	$\max(5,8,6)=8$
B_4	$14 + 6 = 20$	$\max(14,6) = 14$	$\max(14,6) = 14$	$\max(5,8,6,10,12,14) = 14$

Show work below.

Steps:

1. Determine the priority ceiling for each mutex.
2. Under each protocol, for each task, determine the set of lower-priority critical sections that can cause priority inversion to the task: this set consists of the critical sections guarded by mutexes with a priority ceiling greater than or equal to the task priority.

3. Then, apply the properties of the protocol to determine the worst-case blocking time of a task.

Priority Ceilings of Mutexes:

Priority ceiling of M_1 = priority of τ_3

Priority ceiling of M_2 = priority of τ_3

Priority ceiling of M_3 = priority of τ_4

Priority ceiling of M_4 = priority of τ_4

Under PIP:

For τ_4 , the set of lower-priority critical sections to consider: the M_3 critical section of task τ_2 and the M_4 critical section of task τ_3

For τ_3 , the set of lower-priority critical sections to consider: the critical sections of tasks τ_1 and τ_2 that access M_1 and M_2 . Each lower-priority task can contribute only one (outermost) critical section and each mutex can contribute only one critical section. Hence, we have $B_3 = \max(8, 6)$ from τ_2 + 5 from τ_1 .

For τ_2 , the set of lower-priority critical sections to consider: the critical section of M_1 task τ_1 .

Under PCP and HLP: (these have the same blocking times for tasks when tasks do not suspend within a critical section):

For τ_4 , the set of lower-priority critical sections to consider: the M_3 critical section of task τ_2 and the M_4 critical section of task τ_3 = $\max(6, 14)$

For τ_3 , the set of lower-priority critical sections to consider: the critical sections of tasks τ_1 and τ_2 that access M_1 , M_2 and M_3 . Only one of these critical sections can cause priority inversion. Hence, we have $B_3 = \max(5, 8, 6)$.

For τ_2 , the set of lower-priority critical sections to consider: the critical section of M_1 task τ_1 .

For NPP:

Under NPP, any lower-priority critical section can cause priority inversion but at most one can. Hence, B_i is given by the longest critical section of all its lower-priority tasks.

Another equivalent alternative is to consider the priority ceilings of all mutexes to be the highest priority in the system. Hence, all lower-priority critical sections are eligible to cause priority inversion to a task but only one of them can. For the worst case, we pick the longest.